

<https://bit.ly/2I3JKMY>

KYOTO UNIVERSITY

# Statistical Learning Theory - Regression -

Hisashi Kashima

DEPARTMENT OF INTELLIGENCE SCIENCE  
AND TECHNOLOGY

# Linear Regression

# Regression:

## Supervised learning for predicting a real valued variable

- Regression learning is one of supervised learning problem settings with wide applications
- Goal: Obtain a function  $f: \mathcal{X} \rightarrow \mathbb{R}$  ( $\mathbb{R}$  : real value)
  - Usually,  $\mathcal{X}$  is a  $D$ -dimensional vector space
  - E.g.  $x \in \mathcal{X}$  is a house and  $y \in \mathbb{R}$  is its price (housing dataset in UCI Machine Learning Repository)
- Training dataset:  $N$  pairs of an input and an output  $\{(\mathbf{x}^{(1)}, y^{(1)}), (\mathbf{x}^{(2)}, y^{(2)}), \dots, (\mathbf{x}^{(N)}, y^{(N)})\}$

# Some applications of regression:

## From marketing prediction to chemo-informatics

---

- Some applications:
  - Price prediction: Predict the price  $y$  of a product  $x$
  - Demand prediction: Predict the demanded amount  $y$  of a product  $x$
  - Sales prediction: Predict the sales amount  $y$  of a product  $x$
  - Chemical activity: Predict the activity level  $y$  of a compound  $x$
- Other applications:
  - Time series prediction: Predict the value  $y$  at the next time step given the past measurements  $x$
  - Classification

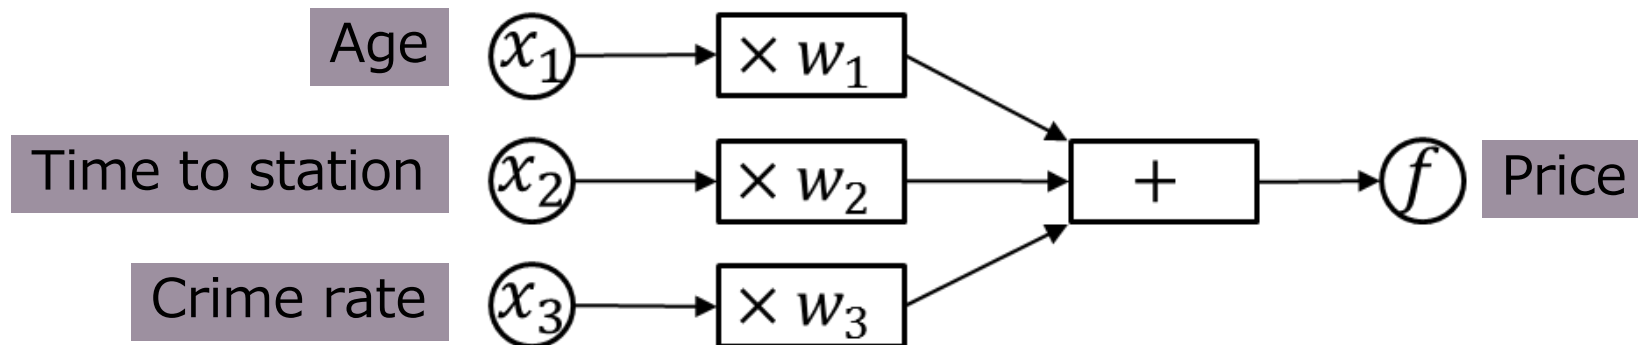
# Model:

## Linear regression model

- Model: How does  $y$  depend on  $\mathbf{x}$ ?
- We consider the simplest choices: Linear regression model

$$y = \mathbf{w}^\top \mathbf{x} = w_1 x_1 + w_2 x_2 + \cdots + w_D x_D$$

—Prediction model of the price of a house:



# Handling discrete features:

## Dummy variables

---

- We assume input  $\mathbf{x}$  is a real vector
  - In the house price prediction example, features can be age, walk time to the nearest station, crime rate in the area, ...
- Discrete features are handled as real values
  - Binary features: {Male, Female} are encoded as {0,1}
  - One-hot encoding: {Kyoto, Osaka, Tokyo} are encoded with (1,0,0), (0,1,0), and (0,0,1)

# Objective function of training:

## Squared loss

---

- Objective function (to minimize):  
Disagreement measure of the model to the training dataset
  - Loss function:  $\ell^{(i)}(y^{(i)}, \mathbf{w}^\top \mathbf{x}^{(i)}; \mathbf{w})$  for the  $i$ -th instance
  - Objective function:  $L(\mathbf{w}) = \sum_{i=1}^N \ell^{(i)}(y^{(i)}, \mathbf{w}^\top \mathbf{x}^{(i)}; \mathbf{w})$

- Squared loss function:

$$\ell^{(i)}(y^{(i)}, \mathbf{w}^\top \mathbf{x}^{(i)}; \mathbf{w}) = (y^{(i)} - \mathbf{w}^\top \mathbf{x}^{(i)})^2$$

- Absolute loss, Huber loss: more robust choices
- Optimal parameter  $\mathbf{w}^*$  is the one that minimizes  $L(\mathbf{w})$ :
$$\mathbf{w}^* = \operatorname{argmin}_{\mathbf{w}} L(\mathbf{w})$$

# Important assumption on data:

## Identically and independently distributed

---

- We assume data are *identically and independently distributed*:
  - Data instances are generated from the same data generation mechanism (or probability distribution)
    - Past data (training data) and future data (test data) have the same property
  - Data instances are independent of each other



# Solution of linear regression:

## One dimensional case

---

- Let us start with a case where inputs and outputs are both one-dimensional
- Objective function to minimize:

$$L(w) = \sum_{i=1}^N (y^{(i)} - wx^{(i)})^2$$

- Solution:  $w^* = \frac{\sum_{i=1}^N y^{(i)} x^{(i)}}{\sum_{i=1}^N x^{(i)2}} = \frac{\text{Cov}(x,y)}{\text{Var}(x)}$

# Solution of linear regression:

## General case

---

- Matrix and vector notations:

- Design matrix  $\mathbf{X} = (\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(N)})^\top$

- Target vector  $\mathbf{y} = (y^{(1)}, y^{(2)}, \dots, y^{(N)})^\top$

- Objective function:

$$\begin{aligned} L(\mathbf{w}) &= \sum_{i=1}^N (y^{(i)} - \mathbf{w}^\top \mathbf{x}^{(i)})^2 = \|\mathbf{y} - \mathbf{X}\mathbf{w}\|_2^2 \\ &= (\mathbf{y} - \mathbf{X}\mathbf{w})^\top (\mathbf{y} - \mathbf{X}\mathbf{w}) \end{aligned}$$

- Solution:  $\mathbf{w}^* = \operatorname{argmin}_{\mathbf{w}} L(\mathbf{w}) = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$

# Regularization

## Ridge regression:

Include penalty on the norm of  $\mathbf{w}$  to avoid instability

- Existence of the solution  $\mathbf{w}^* = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$  requires that  $\mathbf{X}^\top \mathbf{X}$  is non-singular, i.e. full-rank
  - This is often secured when the number of data instances  $N$  is much larger than the number of dimensions  $D$
- Regularization: Adding some constant  $\lambda > 0$  to the diagonals of  $\mathbf{X}^\top \mathbf{X}$  for numerical stability
  - New solution:  $\mathbf{w}^* = (\mathbf{X}^\top \mathbf{X} + \lambda \mathbf{I})^{-1} \mathbf{X}^\top \mathbf{y}$
- Back to its objective function, the solution corresponds to
$$L(\mathbf{w}) = \|\mathbf{y} - \mathbf{X}\mathbf{w}\|_2^2 + \lambda \|\mathbf{w}\|_2^2$$

## Overfitting:

### Degradation of predictive performance for future data

---

- Previously, we introduced the regularization to avoid numerical stability
- Another interpretation: To avoid *overfitting* to the training data
  - Our goal is to make correct predictions for future data, not for the training data
  - Overfitting: Too much adaptation to the training data degrades predictive performance on future data
- When the number of data instances  $N$  is less than the number of dimensions  $D$ , the solution is not unique
  - Infinite number of solutions exist

# Occam's razor:

## Adopt the simplest model

---

- We have infinite number of models that equally fit to the training data (=minimize the loss function)
  - Some perform well, some perform badly
- Which is the “best” model among them?
- Occam's razor: Take the simplest model
  - We will discuss why the simple model is good later in the statistical learning theory
- What is the measure of simplicity?  
For example, number of features = the number of non-zero elements in  $\mathbf{w}$

## 0-norm regularization:

Reduce the number of non-zero elements in  $\mathbf{w}$

- Number of non-zero elements in  $\mathbf{w}$  = 0-norm of  $\mathbf{w}$

- Use 0-norm constraint:

$$\text{minimize}_{\mathbf{w}} \|\mathbf{y} - \mathbf{X}\mathbf{w}\|_2^2 \text{ s. t. } \|\mathbf{w}\|_0 \leq \eta$$

Number of  
features used in  
the model

or 0-norm penalty:

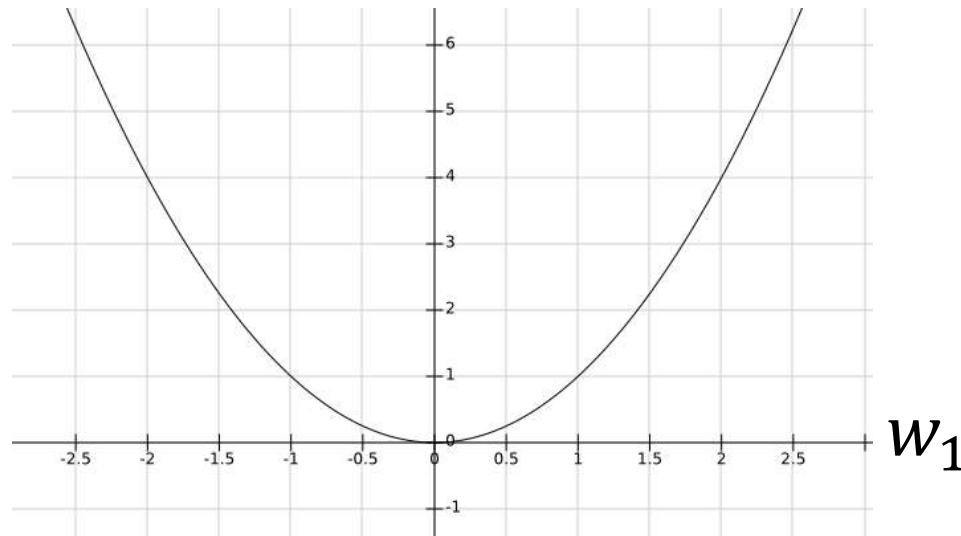
$$\text{minimize}_{\mathbf{w}} \|\mathbf{y} - \mathbf{X}\mathbf{w}\|_2^2 + \lambda \|\mathbf{w}\|_0$$

- There is some one-to-one correspondence between  $\eta$  and  $\lambda$
- However, they are non-convex optimization problems ...
  - Hard to find the optimal solution

# Ridge regression :

## 2-norm regularization as a convex surrogate for 0-norm

- Instead of the zero-norm  $\|\mathbf{w}\|_0$ , we use 2-norm  $\|\mathbf{w}\|_2^2$



Convex 😊

- Ridge regression:  $L(\mathbf{w}) = \|\mathbf{y} - \mathbf{X}\mathbf{w}\|_2^2 + \lambda\|\mathbf{w}\|_2^2$

–Can be seen as a relaxed version of

$$L(\mathbf{w}) = \|\mathbf{y} - \mathbf{X}\mathbf{w}\|_2^2 + \lambda\|\mathbf{w}\|_0$$

Non-convex 😞

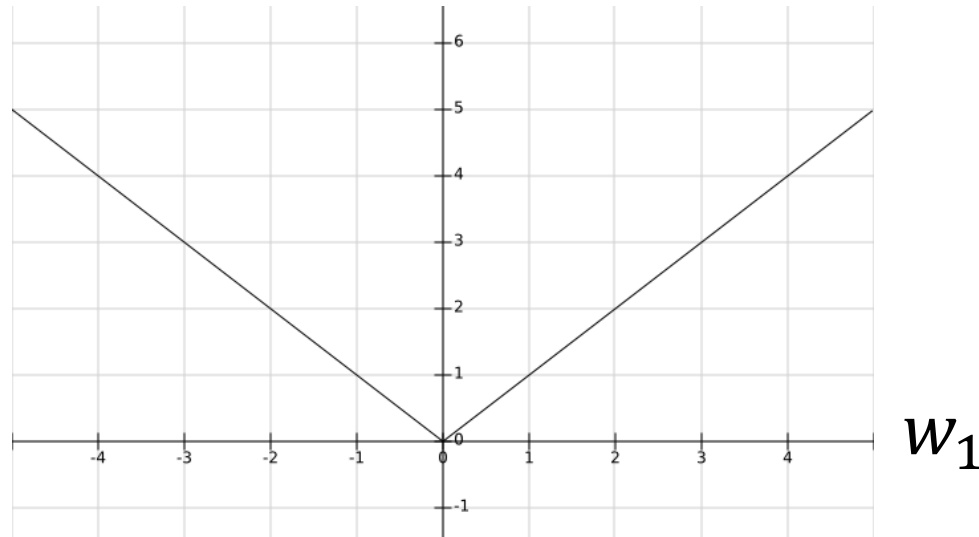
–The closed form solution:  $\mathbf{w}^* = (\mathbf{X}^\top \mathbf{X} + \lambda \mathbf{I})^{-1} \mathbf{X}^\top \mathbf{y}$



## Lasso :

### 1-norm regularization further induces sparsity

- Instead, we can use 1-norm  $\|\mathbf{w}\|_1 = |w_1| + |w_2| + \dots + |w_D|$



- Lasso:  $L(\mathbf{w}) = \|\mathbf{y} - \mathbf{X}\mathbf{w}\|_2^2 + \lambda\|\mathbf{w}\|_1$ 
  - Convex optimization, but no closed form solution
- Sparsity inducing norm: 1-norm induces sparse  $\mathbf{w}^*$

# Statistical Interpretation

# Interpretation as statistical inference :

## Regression as maximum likelihood estimation

---

- So far we have formulated the regression problem in loss minimization framework
  - Function (prediction model)  $f: \mathcal{X} \rightarrow \mathfrak{R}$  is deterministic
  - Least squares: Minimization of the sum of squared losses
- We have not considered any statistical inference
- Actually, we can interpret the previous formulation in a statistical inference framework, namely, maximum likelihood estimation

# Maximum likelihood estimation (MLE):

Find the parameter that best reproduces training data

- We consider  $f$  as a conditional distribution  $f(y|\mathbf{x}, \mathbf{w})$

Conditional probability

- Maximum likelihood estimation (MLE):

- Find  $\mathbf{w}$  that maximizes the likelihood function:

$$L(\mathbf{w}) = \prod_{i=1}^N f(y^{(i)}|\mathbf{x}^{(i)}, \mathbf{w})$$

- Likelihood function: Probability that the training data is reproduced by the model
- Note that we assume i.i.d.

- It is often convenient to use *log* likelihood instead:

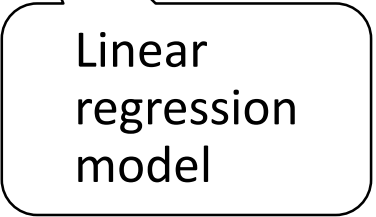
$$L(\mathbf{w}) = \sum_{i=1}^N \log f(y^{(i)}|\mathbf{x}^{(i)}, \mathbf{w})$$

# Probabilistic version of the linear regression model: Gaussian distribution model

---

- Probabilistic version of the linear regression model  $y = \mathbf{w}^\top \mathbf{x}$
- $y \sim \mathcal{N}(\mathbf{w}^\top \mathbf{x}, \sigma^2)$ : Gaussian distribution with mean  $\mathbf{w}^\top \mathbf{x}$  and variance  $\sigma^2$

$$f(y|\mathbf{x}, \mathbf{w}) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(y - \mathbf{w}^\top \mathbf{x})^2}{2\sigma^2}\right)$$



Linear  
regression  
model

# Relation between least squares and MLE:

Maximum likelihood is equivalent to least squares

---

- Log-likelihood function:

$$\begin{aligned} L(\mathbf{w}) &= \sum_{i=1}^N \log f(y^{(i)} | \mathbf{x}^{(i)}, \mathbf{w}) \\ &= \sum_{i=1}^N \log \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(y^{(i)} - \mathbf{w}^\top \mathbf{x}^{(i)})^2}{2\sigma^2}\right) \\ &= -\frac{1}{2\sigma^2} \sum_{i=1}^N (y^{(i)} - \mathbf{w}^\top \mathbf{x}^{(i)})^2 + \text{const.} \end{aligned}$$

- Maximization of  $L(\mathbf{w})$  is equivalent to minimization of the squared loss  $\sum_{i=1}^N (y^{(i)} - \mathbf{w}^\top \mathbf{x}^{(i)})^2$

# Some More Applications

# Time series prediction:

## Auto regressive (AR) model

- Time series data: A sequence of real valued data  $x_1, x_2, \dots, x_t, \dots \in \mathbb{R}$  associated with time stamps  $t = 1, 2, \dots$
- Time series prediction: Given  $x_1, x_2, \dots, x_{t-1}$ , predict  $x_t$
- Auto regressive (AR) model:

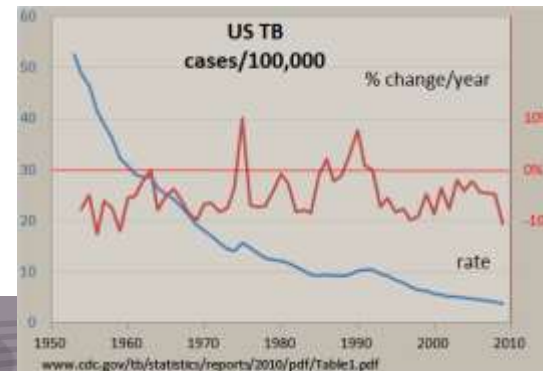
$$x_t = w_1 x_{t-1} + w_2 x_{t-2} + \dots + w_D x_{t-D}$$

—  $x_t$  is determined by the recent length- $D$  history

- AR model as a linear regression model  $y = \mathbf{w}^\top \mathbf{x}$ :

—  $\mathbf{w} = (w_1, w_2, \dots, w_D)^\top$

—  $\mathbf{x} = (x_{t-1}, x_{t-2}, \dots, x_{t-D})^\top$





# Classification as regression:

## Regression is also applicable to classification

---

- Classification:  $y \in \{+1, -1\}$
- Apply regression to predict  $y \in \{+1, -1\}$
- Rigorously, such application is not valid
  - Since an output is either  $+1$  or  $-1$ , the Gaussian noise assumption does not hold
  - However, since solution of regression is often easier than that of classification, this application can be compromise
- Fisher discriminant: Instead of  $\{+1, -1\}$ , use  $\left\{+\frac{1}{N^+}, -\frac{1}{N^-}\right\}$ 
  - $N^+(N^-)$  is the number of positive (negative) data

# Nonlinear Regression

# Nonlinear regression:

## Introducing nonlinearity in linear models

---

- So far we have considered only linear models
- How to introduce non-linearity in the models?
- Introduce nonlinear basis functions:
  - Transformed features: e.g.  $x \rightarrow \log x$
  - Cross terms: e.g.  $x_1, x_2 \rightarrow x_1 x_2$
  - Kernels:  $\mathbf{x} \rightarrow \boldsymbol{\phi}(\mathbf{x})$  (some nonlinear mapping to a high-dimensional space)

# Nonlinear transformation of features:

## Simplest way to introduce nonlinearity in linear models

- Nonlinear basis function:  $x \rightarrow \log x, e^x, x^2, \frac{1}{x}, \dots$

– Sometimes used for converting the range

- E.g.  $\log: \mathbb{R}^+ \rightarrow \mathbb{R}$ ,  $\exp: \mathbb{R} \rightarrow \mathbb{R}^+$

- Interpretations of log transformation:

	$y$	$\log y$
$x$	$y = \beta x + \alpha$ Increase of $x$ by 1 will increase $y$ by $\beta$	$\log y = \beta x + \alpha$ Increase of $x$ by 1 will multiply $y$ by $1 + \beta$
$\log x$	$y = \beta \log x + \alpha$ Doubling $x$ will increase $y$ by $\beta$	$\log y = \beta \log x + \alpha$ Doubling $x$ will multiply $y$ by $1 + \beta$

## Cross terms:

Can include synergetic effects among different features

- Not only the original features  $x_1, x_2, \dots, x_D$ , use their cross terms products  $\{x_d x_{d'}\}_{d,d'}$
- Model has a matrix parameter  $\mathbf{W}$ :

$$y = \text{Trace} \left( \begin{bmatrix} w_{1,1} & \cdots & w_{1,D} \\ \vdots & \ddots & \vdots \\ w_{D,1} & \cdots & w_{D,D} \end{bmatrix}^\top \begin{bmatrix} x_1^2 & x_1 x_2 & \cdots & x_1 x_D \\ x_2 x_1 & x_2^2 & \cdots & x_2 x_D \\ \vdots & \vdots & \ddots & \vdots \\ x_D x_1 & x_D x_2 & \cdots & x_D^2 \end{bmatrix} \right)$$
$$= \mathbf{x}^\top \mathbf{W}^\top \mathbf{x}$$

- $L(\mathbf{W}) = \sum_{i=1}^N \left( y^{(i)} - \mathbf{x}^{(i)\top} \mathbf{W}^\top \mathbf{x}^{(i)} \right)^2 + \lambda \|\mathbf{W}\|_F^2$   
(e.g. factorization machines)

# Kernels:

## Linear model in a high-dimensional feature space

- High dimensional non-linear mapping:  $\mathbf{x} \rightarrow \boldsymbol{\phi}(\mathbf{x})$ 
  - $\boldsymbol{\phi}: \mathcal{R}^D \rightarrow \mathcal{R}^{\bar{D}}$  is some nonlinear mapping from  $D$ -dimensional space to a  $\bar{D}$ -dimensional space ( $D \ll \bar{D}$ )
- Linear model  $y = \bar{\mathbf{w}}^\top \boldsymbol{\phi}(\mathbf{x})$
- Kernel regression model:  $y = \sum_{i=1}^N \alpha^{(i)} k(\mathbf{x}^{(i)}, \mathbf{x})$ 
  - Kernel function  $k(\mathbf{x}^{(i)}, \mathbf{x}) = \langle \boldsymbol{\phi}(\mathbf{x}^{(i)}), \boldsymbol{\phi}(\mathbf{x}) \rangle$ : inner product
  - Kernel trick: Instead of working in the  $\bar{D}$ -dimensional space, we use an equivalent form in an  $N$ -dimensional space
    - Foundation of kernel machines, e.g. SVM, Gaussian process, ...

# Bayesian Statistical Interpretation

# Bayesian interpretation of regression:

## Ridge regression as MAP estimation

---

- Posterior distribution of parameters
- Maximum A Posteriori (MAP) estimation
- Ridge regression as MAP estimation



# Bayesian modeling:

## Posterior, instead of likelihood

---

- In maximum likelihood estimation (MLE), we obtain  $\mathbf{w}$  that maximizes data *likelihood*:

$$P(\mathbf{y} \mid \mathbf{X}, \mathbf{w}) = \prod_{i=1}^N f(y^{(i)} \mid \mathbf{x}^{(i)}, \mathbf{w})$$

$$\text{or } \log P(\mathbf{y} \mid \mathbf{X}, \mathbf{w}) = \sum_{i=1}^N \log f(y^{(i)} \mid \mathbf{x}^{(i)}, \mathbf{w})$$

- The probability of the data reproduced with the parameter:  
 $P(\text{Data} \mid \text{Parameters})$

- In Bayesian modeling, we consider the *posterior distribution*  
 $P(\text{Parameters} \mid \text{Data})$ 
  - Posterior distribution is the distribution over model parameters given data

# Posterior distribution:

Posterior = likelihood + prior

---

## ■ Posterior distribution:

$$P(\text{Parameters} \mid \text{Data}) = \frac{P(\text{Data} \mid \text{Parameters})P(\text{Parameters})}{P(\text{Data})}$$

(Bayes' formula)

## ■ Log posterior:

$$\begin{aligned} \log P(\text{Parameters} \mid \text{Data}) \\ = \underbrace{\log P(\text{Data} \mid \text{Parameters})}_{\text{Likelihood}} + \underbrace{\log P(\text{Parameters})}_{\text{Prior}} \end{aligned}$$

$$-\log P(\text{Data})$$

- $P(\text{Data})$  is often neglected as a constant term because it does not depend on the parameters

# Maximum a posteriori (MAP) estimation:

Find parameter that maximizes the posterior

---

- Log posterior:

$$\log P(\text{Parameters} \mid \text{Data})$$

$$= \log P(\text{Data} \mid \text{Parameters}) + \log P(\text{Parameters}) + \text{const.}$$

- Maximum a posteriori (MAP) estimation finds the parameter that maximizes the posterior:

$$\text{Parameters}^* = \operatorname{argmax}_{\text{Parameters}} \log P(\text{Parameters} \mid \text{Data})$$

- MLE considers only  $\log P(\text{Data} \mid \text{Parameters})$  part

- Additional term (log prior) :  $\log P(\text{Parameters})$

# Ridge regression as MAP estimation:

Find parameter that maximizes the posterior

---

- Log posterior:

$$\log P(\text{Parameters} \mid \text{Data})$$

$$= \log P(\text{Data} \mid \text{Parameters}) + \log P(\text{Parameters}) + \text{const.}$$

- Ridge regression:

$$\mathbf{w}^* = \operatorname{argmin}_{\mathbf{w}} \frac{1}{2\sigma^2} \sum_{i=1}^N (y^{(i)} - \mathbf{w}^\top \mathbf{x}^{(i)})^2 + \frac{1}{2\sigma'^2} \|\mathbf{w}\|_2^2$$

- Log-likelihood:  $\sum_{i=1}^N \log \frac{1}{\sqrt{2\pi}\sigma'} \exp\left(-\frac{(y^{(i)} - \mathbf{w}^\top \mathbf{x}^{(i)})^2}{2\sigma'^2}\right)$

- Prior  $P(\mathbf{w}) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{\mathbf{w}^\top \mathbf{w}}{2\sigma^2}\right)$

# Regression:

## Supervised learning for predicting a real valued variable

- A supervised learning problem to make real-valued predictions
- Regression problem is often formulated as a least-square minimization problem
  - Closed form solution is given
- Regularization framework to avoid overfitting
  - Reduce the number of features: 0-norm, 2-norm (ridge regression), 1-norm (lasso)
- Nonlinear regression
- Statistical interpretations: maximum likelihood estimation, maximum a posteriori (MAP) estimation