

Statistical Learning Theory

- Kernel Methods -

Hisashi Kashima

Kernel Methods

Nonlinear regression:

Introducing nonlinearity in linear models

- In the previous lecture, we saw several ways to introduce non-linearity into linear models by introducing nonlinear basis functions:
 1. Transformed features:
 - e.g. $x \rightarrow \log x$
 2. Cross terms:
 - e.g. $x_1, x_2 \rightarrow x_1 x_2$
 3. Kernel methods:
 - $\mathbf{x} \rightarrow \boldsymbol{\phi}(\mathbf{x})$ (nonlinear mapping to a high-dimensional space)

Dual form of ridge regression:

Parameters as a linear combination of input vectors

- Ridge regression:
 - Prediction model: $y = \mathbf{w}^\top \mathbf{x}$
 - Optimization problem: $L(\mathbf{w}) = \sum_{i=1}^N (y^{(i)} - \mathbf{w}^\top \mathbf{x}^{(i)})^2 + \lambda \|\mathbf{w}\|_2^2$
- Now, we *assume* that the parameter can also be represented as a linear combination of input vectors: $\mathbf{w} = \sum_{i=1}^N \alpha_i \mathbf{x}^{(i)}$
 - Prediction model: $y = \sum_{i=1}^N \alpha_i \mathbf{x}^{(i)\top} \mathbf{x}$
 - Optimization problem:
$$L(\boldsymbol{\alpha}) = \sum_{i=1}^N \left(y^{(i)} - \sum_{j=1}^N \alpha_j \mathbf{x}^{(i)\top} \mathbf{x}^{(j)} \right)^2 + \lambda \sum_{i=1}^N \sum_{j=1}^N \alpha_i \alpha_j \mathbf{x}^{(i)\top} \mathbf{x}^{(j)}$$
 - Note that $\boldsymbol{\alpha} = (\alpha_1, \dots, \alpha_N)$ is the parameter now

Kernel ridge regression:

Ridge regression using kernel function (inner product)

- Observation: inputs are always accessed through **inner product**

- Prediction model: $y = \sum_{i=1}^N \alpha_i \mathbf{x}^{(i)\top} \mathbf{x}$

- Optimization problem:

$$L(\boldsymbol{\alpha}) = \sum_{i=1}^N \left(y^{(i)} - \sum_{j=1}^N \alpha_j \mathbf{x}^{(i)\top} \mathbf{x}^{(j)} \right)^2 + \lambda \sum_{i=1}^N \sum_{j=1}^N \alpha_i \alpha_j \mathbf{x}^{(i)\top} \mathbf{x}^{(j)}$$

- Kernel ridge regression:

- The inner product is called *kernel function*

- Prediction model: $y = \sum_{i=1}^N \alpha_i K(\mathbf{x}^{(i)}, \mathbf{x})$


- Optimization problem:

$$L(\boldsymbol{\alpha}) = \sum_{i=1}^N \left(y^{(i)} - \sum_{j=1}^N \alpha_j K(\mathbf{x}^{(i)}, \mathbf{x}^{(j)}) \right)^2 + \lambda \sum_{i=1}^N \sum_{j=1}^N \alpha_i \alpha_j K(\mathbf{x}^{(i)}, \mathbf{x}^{(j)})$$

Advantage of kernel function:

Introducing non-linearity in linear models

- Consider a (nonlinear) mapping $\phi: \mathbb{R}^D \rightarrow \mathbb{R}^{D'}$
 - D -dimensional space to $D' (\gg D)$ -dimensional space
 - Vector \mathbf{x} is mapped to a high-dimensional vector $\phi(\mathbf{x})$
- Define kernel $K(\mathbf{x}^{(i)}, \mathbf{x}^{(j)}) = \phi(\mathbf{x}^{(i)})^\top \phi(\mathbf{x}^{(j)})$ in the D' -dimensional space
- A linear classifier in the D' -dimensional space is a non-linear classifier in the original D -dimensional space



Advantage of kernel methods:

Computationally efficient (when D' is large)

- Advantage of using kernel function

$$K(\mathbf{x}^{(i)}, \mathbf{x}^{(j)}) = \boldsymbol{\phi}(\mathbf{x}^{(i)})^\top \boldsymbol{\phi}(\mathbf{x}^{(j)})$$

- Usually, computational cost of K should depend on D'
 - D' can be high-dimensional (possibly infinite dimensional)
- But, if we can somehow compute $\boldsymbol{\phi}(\mathbf{x}^{(i)})^\top \boldsymbol{\phi}(\mathbf{x}^{(j)})$ in time depending on D , the dimension of $\boldsymbol{\phi}$ does not matter
- Problem size:
 - D' (number of dimensions) $\rightarrow N$ (number of data)
 - Advantageous when D' is very large or infinite

Example of efficiently computable kernel functions:

Polynomial kernel can consider high-order cross terms

- Combinatorial features: Not only the original features x_1, x_2, \dots, x_D , we use their cross terms (e.g. $x_1 x_2$)
 - If we consider M -th order cross terms, we have $O(D^M)$ terms
- Polynomial kernel: $K(\mathbf{x}^{(i)}, \mathbf{x}^{(j)}) = \left(\mathbf{x}^{(i)\top} \mathbf{x}^{(j)} + c \right)^M$
 - E.g. when $c = 0, M = 2, D = 2$,

$$\mathbf{x}^{(i)} = \begin{pmatrix} x_1^{(i)} \\ x_2^{(i)} \end{pmatrix}$$


$$\begin{aligned} K(\mathbf{x}^{(i)}, \mathbf{x}^{(j)}) &= \left(x_1^{(i)} x_1^{(j)} + x_2^{(i)} x_2^{(j)} \right)^2 \\ &= \left(x_1^{(i)2}, x_2^{(i)2}, \sqrt{2} x_1^{(i)} x_2^{(i)} \right) \left(x_1^{(j)2}, x_2^{(j)2}, \sqrt{2} x_1^{(j)} x_2^{(j)} \right) \end{aligned}$$
 - Note that it can be computed in $O(D)$

Example of efficiently computable kernel functions:

Gaussian kernel has infinite dimensional feature space


- Gaussian kernel (RBF kernel): $K(\mathbf{x}^{(i)}, \mathbf{x}^{(j)}) = \exp\left(-\frac{\|\mathbf{x}^{(i)} - \mathbf{x}^{(j)}\|_2^2}{\sigma}\right)$
 - Can be interpreted as an inner product in an infinite-dimensional space

Discrimination surface with Gaussian kernel



<http://openclassroom.stanford.edu/MainFolder/DocumentPage.php?course=MachineLearning&doc=exercises/ex8/ex8.html>


Gaussian kernel (RBF kernel)



Kernel methods for non-vectorial data:

Kernels for sequences, trees, and graphs

- Kernel methods can handle any kinds of objects (even non-vectorial objects) as long as efficiently computable kernel functions are available
 - Kernels for strings, trees, and graphs, ...



http://www.bic.kyoto-u.ac.jp/coe/img/akutsu_fig_e_02.gif

Representer theorem:

Theoretical underpinning of kernel methods

- Kernel methods rely on the assumption that the parameter is represented as a linear combination of input vectors:

$$\mathbf{w} = \sum_{i=1}^N \alpha_i \mathbf{x}^{(i)}$$

—... but is this theoretically allowed?

- *Representer theorem* guarantees this, when

- Loss ℓ for i -th data depends on \mathbf{w} only through $\mathbf{w}^\top \mathbf{x}^{(i)}$
- L2-regularizer is used

(Simple) proof of representer theorem:

Obj. func. depends only on linear combination of inputs

- Objective function: $L(\mathbf{w}) = \sum_{i=1}^N \ell(\mathbf{w}^\top \mathbf{x}^{(i)}) + \lambda \|\mathbf{w}\|_2^2$
- Divide the optimal parameter \mathbf{w}^* into two parts $\mathbf{w} + \mathbf{w}^\perp$:
 - \mathbf{w} : Linear combination of input data $\{\mathbf{x}^{(i)}\}_i$
 - \mathbf{w}^\perp : Other parts (orthogonal to all input data $\{\mathbf{x}^{(i)}\}$)
- $L(\mathbf{w}^*)$ depends only on \mathbf{w} : $\sum_{i=1}^N \ell(\mathbf{w}^{*\top} \mathbf{x}^{(i)}) + \lambda \|\mathbf{w}^*\|_2^2$
$$= \sum_{i=1}^N \ell \left(\mathbf{w}^\top \mathbf{x}^{(i)} + \underbrace{\mathbf{w}^\perp{}^\top \mathbf{x}^{(i)}}_{=0} \right) + \lambda (\underbrace{\|\mathbf{w}\|_2^2}_{=0} + \underbrace{2\mathbf{w}^\top \mathbf{w}^\perp}_{=0} + \underbrace{\|\mathbf{w}^\perp\|_2^2}_{\text{Minimized to } =0})$$

Kernel logistic regression:

Kernel-based nonlinear classification model

- We can also “kernelize” the logistic regression
- Kernel logistic regression model:

$$f(y = 1|\mathbf{x}) = \frac{1}{1 + \exp(-\mathbf{w}^\top \mathbf{x})} = \frac{1}{1 + \exp\left(-\sum_{i=1}^N \alpha_i K(\mathbf{x}^{(i)}, \mathbf{x})\right)}$$

- Objective function of kernel (regularized) logistic regression:

$$L(\boldsymbol{\alpha}) = \sum_{i=1}^N \ln\left(1 + \exp\left(-y^{(i)} \sum_{j=1}^N \alpha_j K(\mathbf{x}^{(i)}, \mathbf{x}^{(j)})\right)\right) \\ + \lambda \sum_{i=1}^N \sum_{j=1}^N \alpha_i \alpha_j K(\mathbf{x}^{(i)}, \mathbf{x}^{(j)})$$